

How Operating Systems Work, Why You Should Build Your Own, and How to Do It

by Lukas Yoder



My Pet Project: DragonOS

If you want, you can clone the repo and build the OS during this presentation.

Link: <https://github.com/SentToDevNull/DragonOS>

A screenshot of a GitHub repository page for 'DragonOS' owned by 'SentToDevNull'. The page features a header with navigation links for Code, Issues (0), Pull requests (0), Projects (0), Wiki, Insights, and Settings. Below the header, a brief description states: 'A source-based and modular build system for DragonOS, a minimalist system based on the Linux kernel.' Key statistics shown include 6 commits, 1 branch, 0 releases, and 1 contributor. A green horizontal bar spans the width of the stats. Below these, there are buttons for Branch: master, New pull request, Create new file, Upload files, Find file, and Clone or download. The main content area displays a list of recent commits:

Commit	Message	Time
SentToDevNull	Upgraded base system components and switched to universal extract scr...	Latest commit 1b0d844 14 hours ago
parts	Upgraded base system components and switched to universal extract scr...	14 hours ago
README.md	Added README with TODO info	14 hours ago
dragonos-builder.sh	Bumped version and added support for universal extract script.	14 hours ago

At the bottom of the page, a large grey banner contains the text 'DragonOS' in bold black font. To the left of this banner is a small red square containing the number '2'.

How I Got Started with Working on This



The Problem

DontBreakDebian

Advice For New Users On Not Breaking Their Debian System

Debian is a robust and reliable system, but it's still very easy for new users to break their systems by not doing things the Debian way. This page lists common mistakes made by new users. Some of the things listed here can be done safely, but only if you have enough experience to know how to fix your system when things go wrong.

The general theme to the advice here is that consequences are not always immediate, and can make future upgrades impossible without a complete reinstall. If upgrading without a complete reinstall is important to you, be careful not to make the mistakes outlined below.

One of the primary advantages of Debian is its central repository with thousands of software packages. If you're coming to Debian from another operating system, you might be used to installing software that you find on random websites. On Debian **installing software from random websites is a [bad habit](#)**. It's always better to use software from the official Debian repositories if at all possible. The packages in the Debian repositories are known to work well and install properly. Only using software from the Debian repositories is also much safer than installing from random websites which could bundle malware and other security risks.



Don't make a FrankenDebian

Debian Stable should not be combined with other releases. If you're trying to install software that isn't available in the current Debian Stable release, it's not a good idea to add repositories for other Debian releases. The problems might not happen right away, but the next time you install updates.

The reason things can break is because the software packaged for one Debian release is built to be compatible with the rest of the software for that release. For example, installing packages from *buster* on a *stretch* system could also install newer versions of core libraries including  [libc6](#). This results in a system that is not *testing* or *stable* but a broken mix of the two.

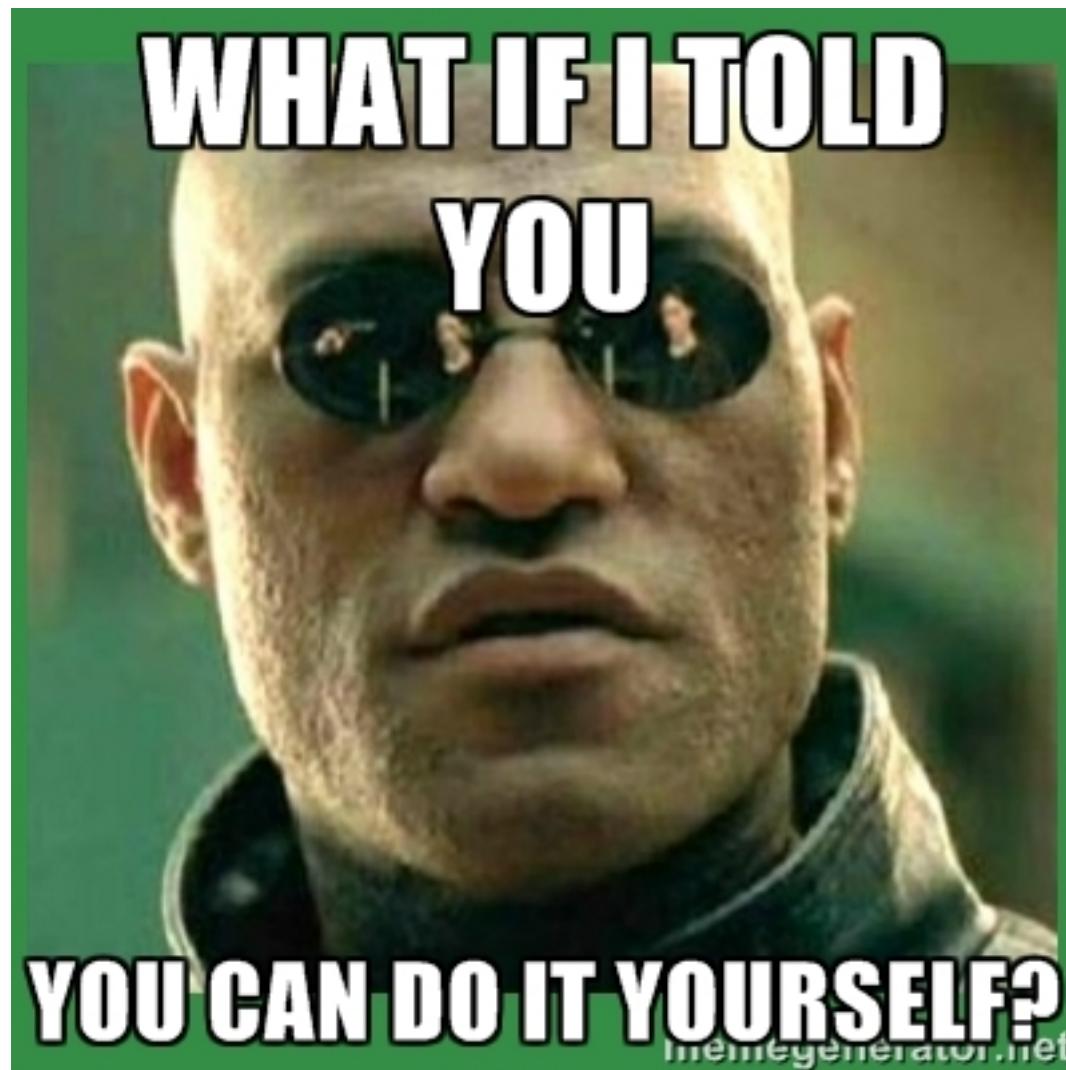
Repositories that can create a FrankenDebian if used with Debian Stable:

- Debian *testing* release (currently *buster*)
- Debian *unstable* release (also known as *sid*)
- Ubuntu, Mint or other derivative repositories are **not** compatible with Debian! |
- Ubuntu PPAs

The Problem, Cont.



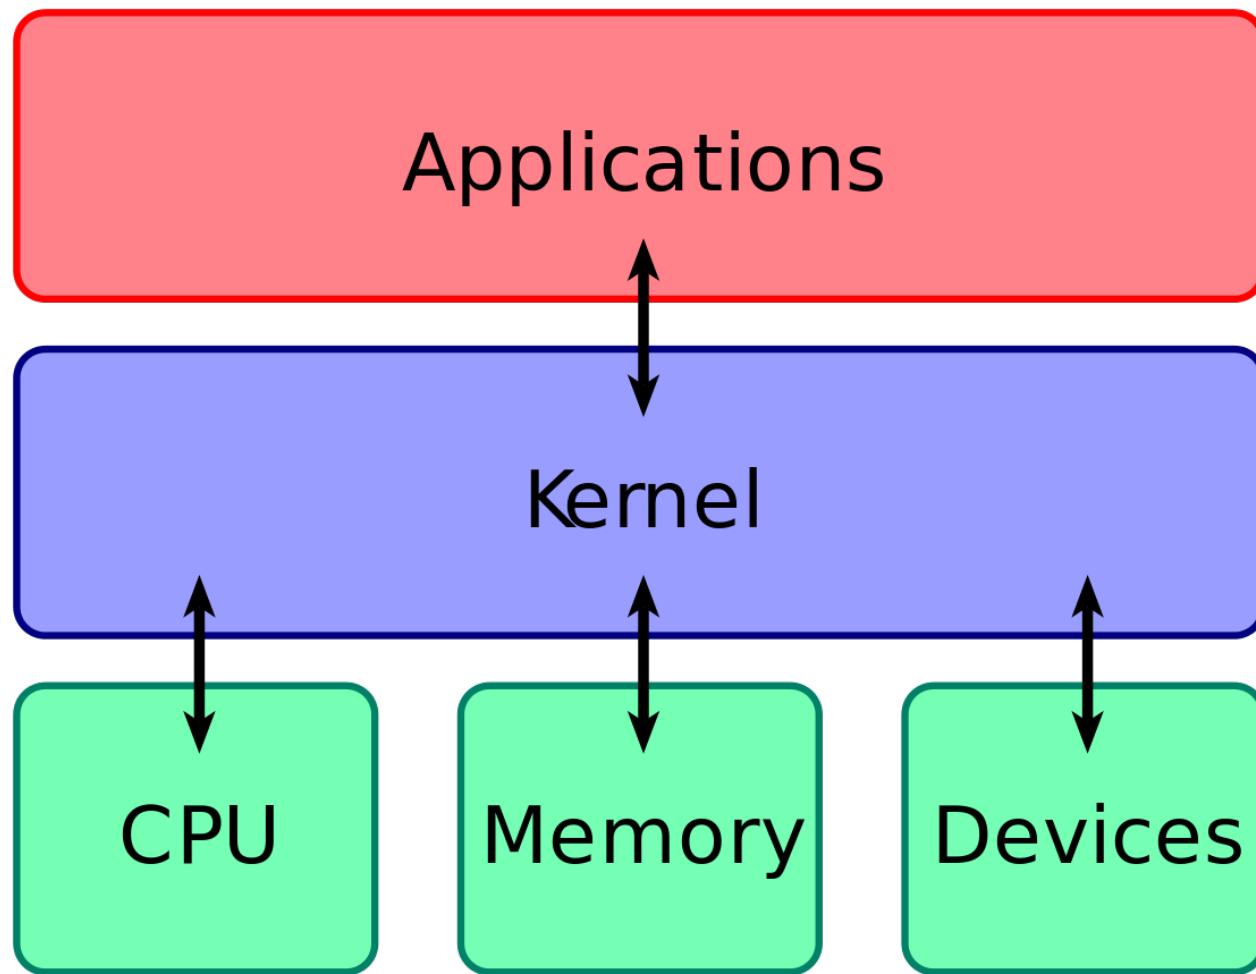
Solution: Do it Yourself



...



How Operating Systems Work: Part 1 – The Kernel



See How Kernels Work: A Hello World Kernel

kernel.asm

```
;architecture
bits 64

;virtual address space
section .text
;multiboot spec mandates the following defined within the first 8kb:

align 4 ;have a 4-byte boundary for instructions

;a magic number to identify the header (occupying first 4 bytes)
; using "dd" for double-word (4 bytes)
dd 0x1BADB002;0x1BADB002 is the magic number for a multiboot header

;a "flags" field occupying the next 4 bytes
;if bit 0 is set, all boot modules are loaded aligned on page boundaries
;if bit 1 is set, info on available memory must be included
;if bit 2 is set, info about video mode table must be included
dd 0x00000000 ;setting bit 0

;a "checksum" field containing a uint32 that when added to "magic" and "flags" must equal 0
dd - (0x1BADB002 + 0x00000000)

;end of multiboot section

extern function ;the function called "function" that is defined externally...

global start ;need to declare the start segment globally so the linker can find it
start:
    mov esp, this_is_your_stack ;move the extended stack pointer to your stack
    call function ;self-explanatory
    hlt ;let's prevent a stack overflow; if I were decent at assembly, this wouldn't be necessary

section .bss ;your friendly neighborhood data segment for statically-allocated variables
this_is_your_stack: resb 8192 ;reserving 8192 bytes for the stack
```

kernel.c

```
void function(void)
{
    const char *str = "Heya there, world!";
    //video memory begins at address 0xb8000
    char *vidptr = (char*)0xb8000;
    unsigned int i = 0;
    unsigned int j = 0;
    unsigned int screensize;

    //this loops clears the screen
    //there are 25 lines each of 80 columns; each element takes 2 bytes
    screensize = 80 * 25 * 2;
    while (j < screensize) {
        //blank character
        vidptr[j] = ' ';
        //attribute-byte
        vidptr[j+1] = 0x07;
        j = j + 2;
    }

    j = 0;

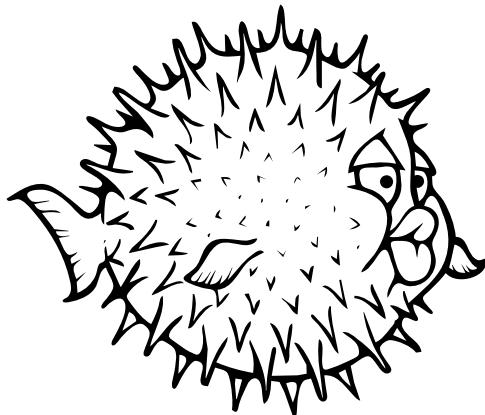
    //this loop writes the string to video memory
    while (str[j] != '\0') {
        //the character's ascii
        vidptr[i] = str[j];
        //attribute-byte: give character black bg and light grey fg
        vidptr[i+1] = 0x07;
        ++j;
        i = i + 2;
    }

    return;
}
```

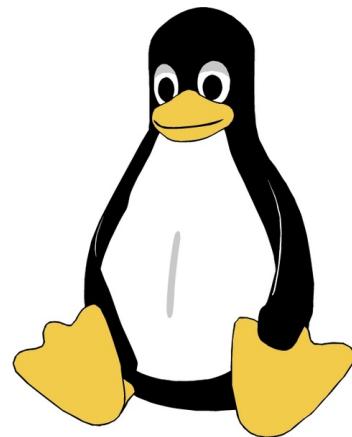
A Buffet of Kernels



Strength: Most Widespread; on Every Motherboard with Intel Chipset; Runs the Management Engine



Strength: Most Paranoidly Secure; Not the Best Performance



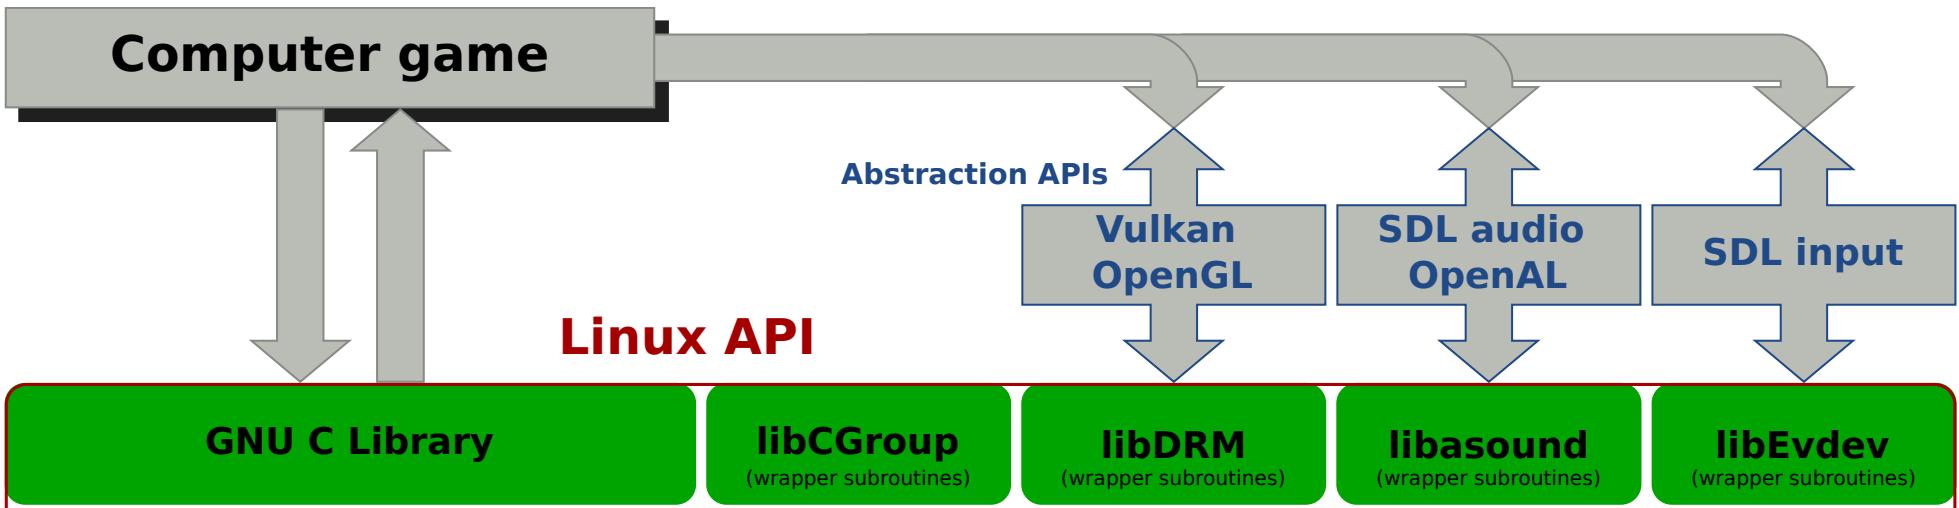
Strength: Most Hardware Support; Good Middle Ground b/w Performance & Security



Strength: Best Performance; Least Security

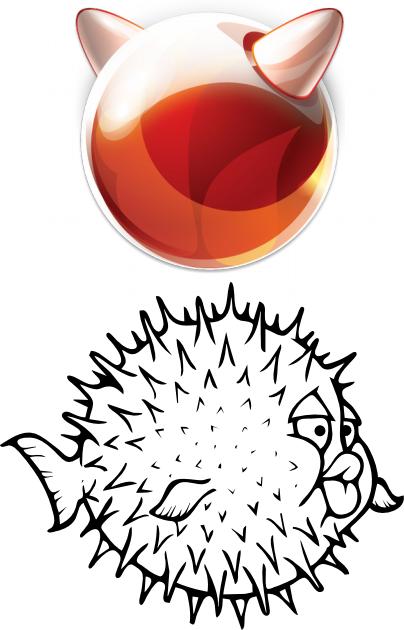


Climbing up the Ladder of Abstraction: C Libraries

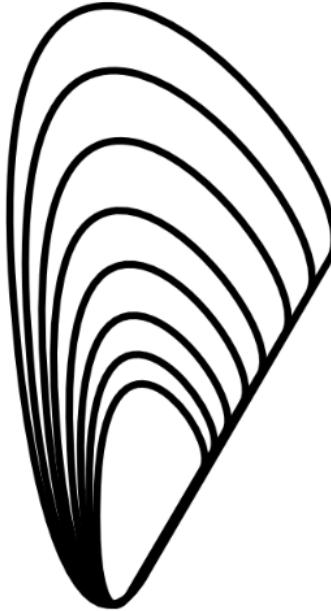


Hardware

Can you C all These Different Libraries?



BSD Libc: Each BSD maintains their own; Android uses a fork
OpenBSD's called Bionic



musl libc: a lightweight, fast, simple, and free libc implementation; does pretty much everything glibc does with less code, less vulnerabilities, and better performance



Glibc: The GNU project's libc implementation; has been around forever; most widely used; most features; also most bloated :'



Climbing Higher Up the Ladder of Abstraction: System Utilities



Most Important Utility: Your CLI Shell; it's a language and an interface to all programs you have installed on your system; alternatives: ZSH, KSH, CSH, DASH

```
89 -----#
90 #                                     #
91 #      Set up a bootstrapped environment to build the system... #
92 #                                     #
93 #-----#
94
95 function BOOTSTRAPSETUP {
96 cp parts/04-bootstrap-setup.sh $LFS
97 chown -v lfs /mnt/lfs/tools
98 chown -v lfs /mnt/lfs/sources
99 cp parts/extract.sh $LFS/sources/
100 time su lfs -c /bin/bash << "EOF"
101 exec env HOME=$HOME TERM=$TERM PS1='\u:\w\$ ' LFS=/mnt/lfs LC_ALL=POSIX \
102     LFS_TGT=$(uname -m)-lfs-linux-gnu PATH=/tools/bin:/bin:/usr/bin bash \
103     /mnt/lfs/04-bootstrap-setup.sh
104 EOF
105 rm -f $LFS/04-bootstrap-setup.sh
106 }
```

... and GitHub messed
up the coloring in its
code viewer...

Climbing Higher Up the Ladder of Abstraction: More System Utilities

zlib: a data compression library and utility

file: determines the filetype of a program

readline: line editing library that allows you to
do things like move your text cursor

M4: a macro processor

bc: a “basic calculator”

binutils: programming tools for creating and managing
binary programs, object files, libraries,
profile data, and assembly code

GMP: a multiple precision arithmetic library

MPFR: an arbitrary-precision binary floating-point
computation library with correct rounding

MPC: a library for the arithmetic of complex numbers

...

ncurses: a text-based user interface library

bzip2: another compression algorithm

GCC: the GNU compiler collection, which contains compilers
for C, C++, Objective-C, Objective-C++, Fortran, Java,
Ada, Go, etc.

sed: a stream editor; used a lot for text manipulation

perl: a high-level interpreted language

coreutils: software that implements many basic unix tools

util-linux: software implementing over 100 basic linux
system utilities not in coreutils

GRUB: GNU’s Grand Unified Bootloader

man: a utility that helps you browse manual pages

vim: the one true text editor

... and many more...

Note 1: A lot of these tools don’t just have 1 implementation, but rather many different ones. For example, GNU Binutils is not the only binutils implementation. Some of these utilities can be substituted for others that perform similarly, but a functional system requires all of these.

Note 2: Some of the utilities here are deemed necessary by the POSIX (Portable Operating System Interface) standard.

Note 3: These aren’t necessarily the names of command-line utilities, but rather pieces of software that may consist of more than 1 program.

Why is all this necessary?

Abstraction like this is necessary so that you don't have to worry about buffers, registers, and CPU microarchitectures to do basic tasks.

```
[user@system ~]$ echo 1+2 | bc  
3
```

Climbing Higher Up the Ladder of Abstraction: Init Systems

Init systems are the first processes started when operating systems boot. They manage other system processes. (For instance, your network service that connects you to WiFi after you turn your computer on.)

There are many different init systems to choose from, including:

- Systemd
- Launchd
- System V
- Upstart

Different init systems have different strengths and weaknesses. For example, systemd is very fast compared to System V, which used to have more widespread use, but System V is preferred by some because it follows the Unix philosophy by not being monolithic in design.

Climbing Higher Up the Ladder of Abstraction: Display Servers



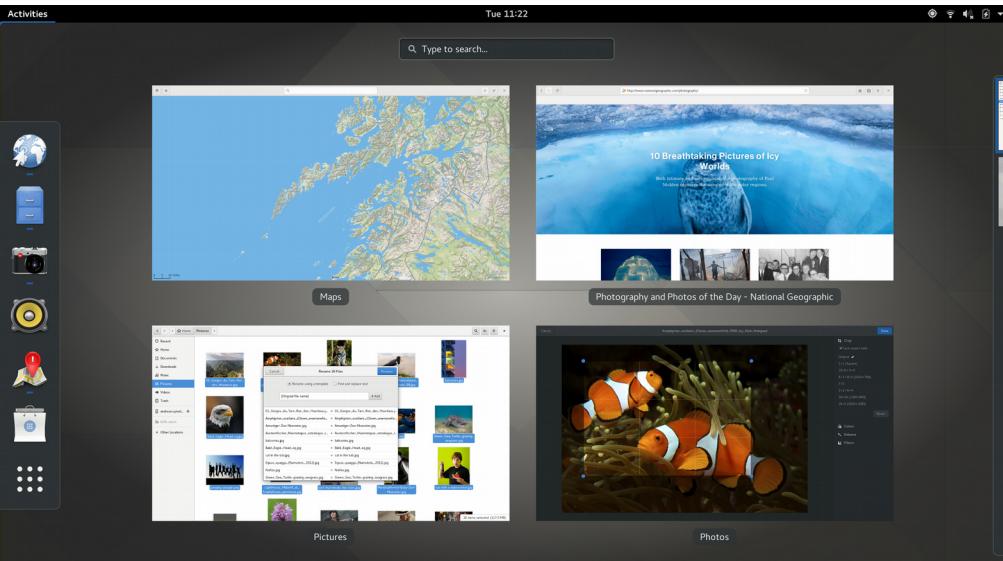
X.Org



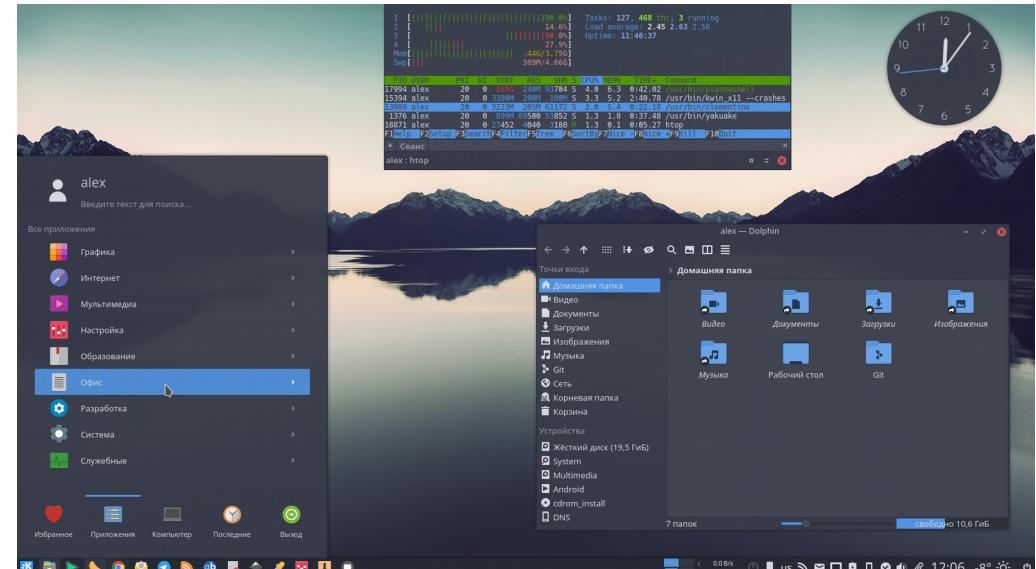
Wayland; the successor
of the X.Org project, intended
to be more secure, among
other things

Climbing Higher Up the Ladder of Abstraction: Graphical Shells

Sometimes you don't need the power of a CLI shell and it's just easier to point and click.



GNOME 3



KDE Plasma 5

Building Your Own System

Not easy... Occasionally not very fun...

Some Challenges:

- Software components don't always play nice with each other.
 - ex: the more recent version of util-linux doesn't seem to compile with ncurses, an optional dependency
- Sometimes base system software requires patches to run on your platform. (Sometimes it needs patches to function at all.)

Something like the Linux From Scratch Project can provide a nice base for your operating system, since it records what software is compatible with other software and what needs to be done to make certain software compile and work with your system.

Linux From Scratch Project

If you're interested in operating system development, the LFS manual is a great way to learn about ONE way to do it.

If you want to make a system that you can actually maintain and add to, the LFS project's instructions aren't the best way to go.

- Take the basic concepts and implement them the way you feel is best.
- The way LFS structures the steps in their manual is to maximize educational value, not to be the most efficient. (That's even stated somewhere in the documentation, if you can manage to find it.)
 - For example, the way that the LFS project is structured prevents automation through the shell. (Indeed, their automated LFS project has to use a Java program to parse through the PDF textbook and use system libraries to build their OS.)

Things That Are Probably Useful to Implement

- A package manager to keep track of the software you build.
- Automated logging to enable you to resume builds if they error out at any point (or if the power to your PC shuts off).
- Automation of everything that the LFS manual wants you to do manually, such as checking software on the host system.
- Writing to an ISO file instead of an actual partition.
- A build system that enables you to easily debug problems with your build and restart or continue it if necessary.
- A package cache so that you don't have to keep re-downloading files.
- Any changes you feel make the system more awesome.

... and that's where my pet project comes into play

Among other things, it adds the stuff discussed in the last slide to make development easier.

If you're interested in building your own unique system, you can read the LFS manual, read these slides, and get started hacking.

If you just want to test a system like this out, you can clone my repo (<https://github.com/SentToDevNull/DragonOS>), run the build script, install (development versions of) dependencies, and have fun.

Now, A Demo of DragonOS